

Runtimes & Execution Engines



UNIVERSITY OF
COPENHAGEN



I Three families of systems

Classic Microservices

Spring, Flask, **Dapr**, etc.

Database + Serverless Functions

Boki, Beldi, Cloudburst, DBOS

Dataflow-based

Styx

Statefun

Actor-like

Orleans

Microservices

Orchestrations, service mesh, databases, etc.

Spring Framework

Java-based web application framework

Standardized interfaces to facilitate operational concerns

Security


Database API

Network

Modularity

Testing

Industry-strength framework for cloud applications



Source: Spring Framework Docs (2025)

I Spring through a “data management” perspective


ORM-heavy specifications

Developers specify local transactions via annotations
Limited support for distributed transactions

Not popularly pursued

Wide library support for event-driven communication


Safety and liveness at developers' hands




Source: Spring Framework Docs (2025)

Dapr

- Distributed framework for event-driven applications
- Standardized APIs that abstract away technologies
- Side-car pattern
- Configurable message delivery
- Workflow management




A bird's eye view into Dapr



Dapr through a “data management” perspective


- Performance overhead
- Centralized state management
- No relational model
- No strong transactional guarantees
- Weak delivery semantics



Actors

Akka, Orleans, etc.

Orleans: Invoking a method on actor A



Source: Philip A. Bernstein, Sergey Bykov:
*Developing Cloud Services Using the
Orleans Virtual Actor Model. IEEE Internet
Comput. 20(5): 71-75 (2016)*

Overview of Orleans Runtime


Messaging multiplexed over a few TCP connections

Actors run on a small number of threads, one per core

Actor Directory is a custom DHT

Cluster Membership management relies on a reliable distributed database


Source: Philip A. Bernstein, Sergey Bykov:
Developing Cloud Services Using the Orleans Virtual Actor Model. IEEE Internet Comput. 20(5): 71-75 (2016)



Actor State Management

The runtime instantiates an actor and reads the actor's state from storage

```
public CartActor([PersistentState(  
    stateName: "cart",  
    storageName: Constants.OrleansStorage)] IPersistentState<Cart> state)  
{  
    this.cart = state;  
}
```



Transactions in the Actor Model

Many applications require atomic operations, possibly over multiple actors, e.g.,

A checkout workflow involving Cart, Stock, Payment and Shipment actors

Transaction management on actor abstraction is challenging.

Multi-actor transactions are distributed transactions, even if actors are collocated on a single server

Orleans Transaction

2PL to achieve concurrent control

Commit using 2PC

Early lock release:

- Releasing locks during phase one of 2PC

- Reducing excessive latency caused by locking of data on cloud storage

Deadlock can be avoided by sorting the access order of actors

- Use timeout if sorting is not possible

I Snapper: A transaction library on actor systems

Transactions can be executed in two modes

Deterministic: **PACT** (Pre-declared ACtor Transaction)

Non-deterministic: **ACT** (ACtor Transaction)

PACT: Trans are ordered and executed deterministically in each involved actor

App code declares the set of actors involved in a transaction

ACT: 2PL + 2PC + no wait

Novel **hybrid** execution mechanism:

Snapper allows both PACT and ACT to run simultaneously

A global serializability check to ensure correctness

Source: *Yijian Liu, Li Su, Vivek Shah, Yongluan Zhou, Marcos Antonio Vaz Salles:*

Hybrid Deterministic and Nondeterministic Execution of Transactions in Actor Systems. SIGMOD Conference 2022: 65-78

Snapper: A transaction library on actor systems


Smallbank Benchmark

Contention \uparrow

Orleans Txn \downarrow

ACT \downarrow



PACT \uparrow



Source: *Yijian Liu, Li Su, Vivek Shah, Yongluan Zhou, Marcos Antonio Vaz Salles: Hybrid Deterministic and Nondeterministic Execution of Transactions in Actor Systems. SIGMOD Conference 2022: 65-78*

Snapper: A transaction library on actor systems

Hybrid execution:




Source: Yijian Liu, Li Su, Vivek Shah, Yongluan Zhou, Marcos Antonio Vaz Salles:
Hybrid Deterministic and Nondeterministic Execution of Transactions in Actor Systems. SIGMOD Conference 2022: 65-78

Serverless Functions

Lambda/Azure Functions, Durable Functions, Amazon Steps

Function as a Service – Typical System



Serverless Function as a Service

Transparent application management

Users rely that application functions are:

- persistent (no submission of the function for every request)

- scheduled (in a timely and fair manner)

- instantiated (provision of resources and execution)

- scaled (according to application workload)

- persisted (results are stored durably)


Serverless Function as a Service - History

Lambda paradigm (e.g., AWS Lambda)

Small to moderate I/O (storage
and network) workloads

Lack of execution guarantees
intra and inter functions

Costly shared and mutable state




Moellering & Grunwald in AWS Architecture Blog (2020)

Serverless Function as a Service - History

Stateful FaaS (SFaaS) (1st gen)

Specific-purpose
autoscaling storage
Richer data consistency
semantics
Weak guarantees on
multi-function workflows



I Serverless Function as a Service - History

SFaaS (2nd generation)
Actor-oriented
programming model
Stronger execution
guarantees




Stateful
Functions

KALIX



Serverless Function as a Service – Typical System




Serverless Function as a Service

Function execution is multiplexed in a virtual layer

Virtualized sandboxes

1. Function is registered
2. Function is invoked
3. Upstream layers process invocation
4. Virtualized sandbox is loaded/reused
5. Sandbox executes function in isolation



Serverless Function as a Service

Performance techniques

Reuse a warm sandbox

Avoid cold start

How long to keep?

Pre-create a sandbox

Overprovision

predicting load increase

Balance load

Increase latency

Distribute and batch functions

according to previous runs

Serverless Function as a Service

FaaS is a model for executing functions

FaaS system abstracts operational concerns to enable FaaS model

- Resource and failure transparency to developer


- Focus on writing application functions

Complex, data-intensive applications may not fit well the paradigm (yet)

Application architecture is always a trade-off

Dataflow systems

| Styx



Object-oriented API

Dataflows are awkward.
Styx offers Durable Entities [4],
i.e., Python objects with
arbitrary function-to-function
calls.

"Keep the data moving" [1]

No time for wait for 2PCs. Styx
extends deterministic database
concepts [2] for arbitrary
function-to-function calls.

Scalability

Partitioned state, collocated
with function execution.
Parallel execution at Entity-
level granularity.

Coarse-grained Fault Tolerance

Async checkpoints á lá
Chandy-Lamport [3,4].

Exactly-once output

Uses durable queues (Kafka) for
input transactions, and
deduplicating outputs.

Consistency

Serializable state mutations
made by arbitrary function-to-
function calls.

[1] Stonebraker, Çetintemel, Zdonik. "The 8 requirements of real-time stream processing." [Sig. Record 2005]

[2] Yi, Yu, Cao, and Samuel Madden. "Aria: a fast and practical deterministic OLTP database." [VLDB 2020]


[3] Carbone, et.al. "State management in Apache Flink®: consistent stateful distributed stream processing." [VLDB 2018]

[4] Silvestre, Fragkoulis, Katsifodimos. "Clonos: Consistent causal recovery for highly-available streaming dataflows." [SIGMOD 21]

[5] Psarakis, Zorgdrager, Fragkoulis, Salvaneschi, Katsifodimos "Stateful Entities: Object-oriented Cloud Applications as Distributed Dataflows", [CIDR '23, EDBT '24]

[6] Psarakis, Christodoulou, Siachamis, Fragkoulis, Katsifodimos "Styx: Transactional Stateful Functions on Streaming Dataflows", [SIGMOD'25]


Enter Styx: an “ideal” Transactional Cloud Application Runtime




<https://github.com/delftdata/styx>

34

Stateful Entities (WIP): Object-oriented Style Programming



Deployment on Styx



Python	Dataflow
Class =>	Operator
Object State =>	Operator State
Function Call Arguments =>	Event
Return Value =>	Event


Stateful Dataflow Graph

- Parallelizable
- Exactly-once processing guarantees
- Transactional serializability
- High-throughput, low-latency

Event-driven programming in Styx (Operator API)

```
@cart.register
def checkout(ctx: StatefulFunction):
    items, user_id, total_price, paid = ctx.get()
    for item_id, qty in items:
        ctx.call_async(operator=stock,
                      function_name='decrement_stock',
                      key=item_id,
                      params=(qty, ))
    ctx.call_async(operator=payment,
                  function_name='pay',
                  key=user_id,
                  params=(total_price, ))
    paid = True
    ctx.put((items, user_id, total_price, paid))
    return "Checkout Successful"
```


How does the architecture look like?




I Styx is epoch-based

How does it perform?

DeathStar Throughput vs. Latency





YCSB Throughput vs. Latency



100 workers/CPUs

How does it perform?



Summary

Existing FaaS avoid important problems

State, messaging & transactions need a holistic solution

It is possible to build “ideal” runtimes

